



Deep Generative Model: A Statistical Perspective

Young-geun Kim

Department of Statistics and Probability

STT 990 (Fall 2024)

Outline

1 INTRODUCTION

2 STATISTICAL DISTANCES IN DEEP GENERATIVE MODELS

- f Divergence-based Methods
- Integral Probability Metric-based Methods
- Wasserstein Distance-based Methods
- Fisher Divergence-based Methods

What are Generative Models?

- The term ‘Generative Model’ has been used in classification since before the emergence of deep generative models.
- Let X and Y represent observations and labels, respectively, in a classification problem.
 - ① Generative Model: Models $p(X, Y)$, typically by modeling $p(X|Y)$ and $p(Y)$ separately.
 - ② Discriminative Model: Models $p(Y|X)$ only and uses it directly.
- Since the generative model learns $p(X, Y)$, it can generate data, e.g., by first sampling $Y \sim \text{Ber}(p)$ and then sampling $X|Y \sim N(\mu_Y, \sigma_Y^2)$ to synthesize (X, Y) pairs.

What are Generative Models?

- The advent of high-dimensional and large-scale data has increased the richness of information within the data, even without human-annotated labels, emphasizing the need for advanced statistical methods.
- Recent advancements have enabled the learning and generation of complex data. These models are now commonly referred to as Generative AI or Generative Models.
- Similar to traditional generative models used in classification, they learn the joint distribution of all observed variables, \vec{X} :

$$p(\vec{X}). \quad (1)$$

This capability allows them to generate new data samples.

What are Generative Models?

- In this talk, the term ‘Generative Model’ refers to statistical models that learn the distribution of observations either without human-annotated labels or with auxiliary information.¹
- Deep generative models have shown remarkable performance as:
 - ① simulators by generating realistic data,
 - ② dimension reduction tools by extracting low-dimensional representations,
 - ③ inference tools by translating observations to other domains.

¹For example, demographic information in medical data or timestamps in temporal data.

Application: Image Generation



- After training, generative models can synthesize realistic data without prior information such as age and race in face generation.
- These models can be used to augment data, generate privacy-free samples, and enhance virtual reality experiences.

The generated image is from <http://thispersondoesnotexist.com/>

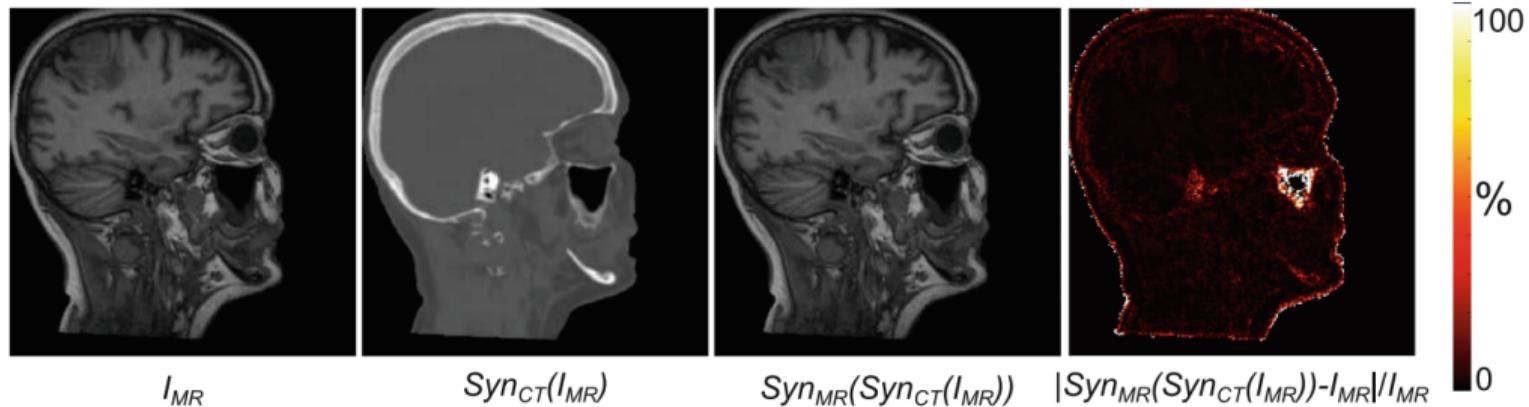
Application: Text-to-Image Generation



- Images are generated by DALL-E-2 using
*"There is a clean desk in the middle. Outside the window,
a whale shark is swimming in the dark night sky above Manhattan."*
- Generated images reflect semantic information in text descriptions.

Images are generated with DALL-E-2 (Ramesh et al., 2022).

Application: Cross Modality Transfer



- Generative models are also useful in imputing missing modalities, e.g., $MR \rightarrow CT$, or enhancing data resolution.

Images are from Wolterink et al. (2017).

Challenges in Deep Generative Models



- **High-dimensional Data:** For data like 4K-resolution color images, the dimensionality is $3,840 \times 2,160 \times 3 (\approx 24M)$.
- **Complex Structure:** Image, video, audio, and language data exhibit complex structures.

Challenge in Deep Generative Models

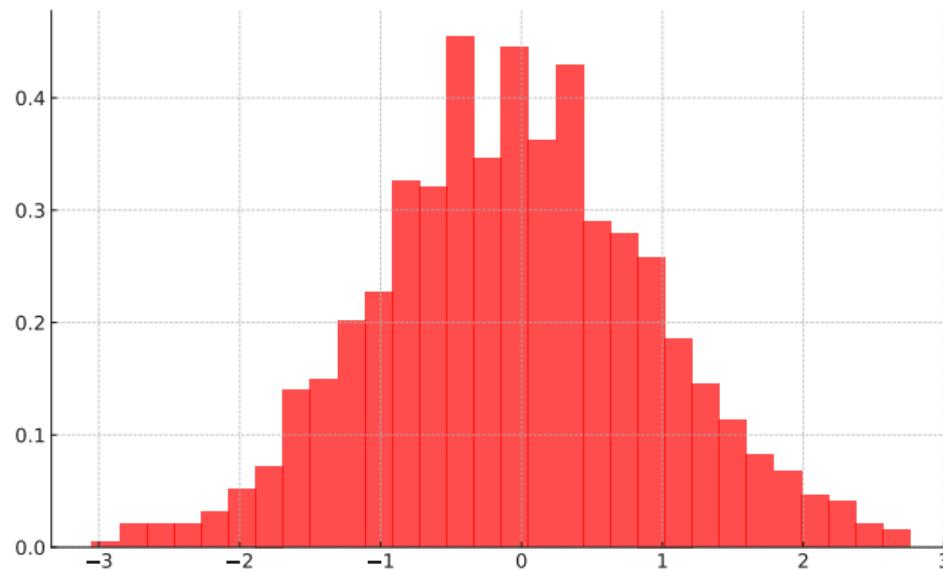


- How to model the distribution of high-dimensional data efficiently?
- How to evaluate the generated data and train the model distribution?

Example of generated images, edited from Li et al. (2024)

How to Learn Distributions: Toy Example

- Let's begin with a basic example. Assume we observed n univariate samples x_1, \dots, x_n having the following histogram:



How to Learn Distributions: Toy Example

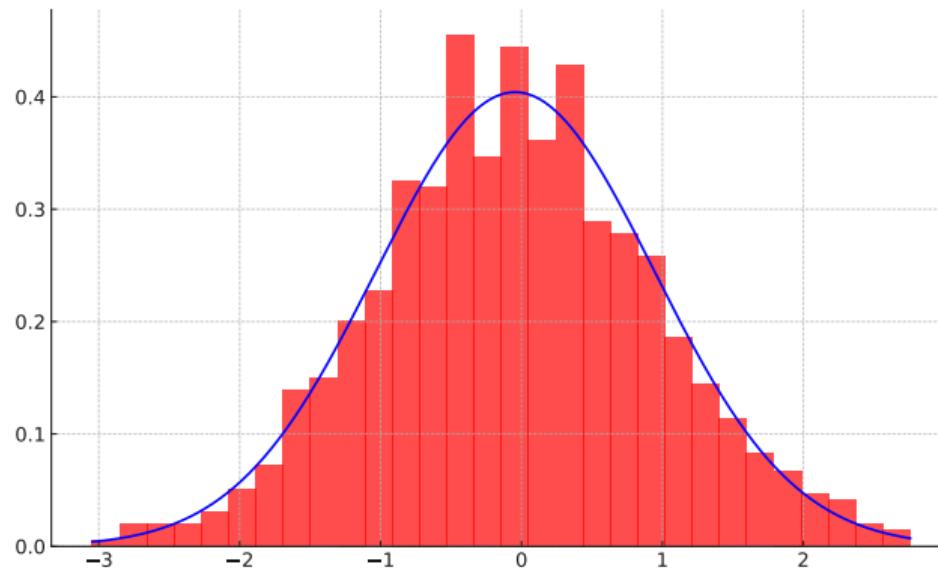
- How to learn the distribution where x_i comes from?
- One way is to model it as Gaussian and find the optimal mean and std parameters.
- $p_\theta(x) := (2\pi\sigma^2)^{-1/2} \exp\left((x - \mu)^2/(2\sigma^2)\right)$ where $\theta = (\mu, \sigma^2)^T \in \mathbb{R} \times \mathbb{R}^+.$ ²

Q: What are good evaluation criteria for optimality? How do we determine which parameter values are superior?

²For brevity, parameter vectors are denoted without vector symbols if there is no confusion.

How to Learn Distributions: Toy Example

- $I_n(\theta) := \sum_{i=1}^n \log p_\theta(x_i) = -(n/2) \log \sigma^2 - \sum_{i=1}^n (x_i - \mu)^2 / (2\sigma^2) - (n/2) \log 2\pi$
- $\arg \max_{\theta} I_n(\theta) = (\hat{\mu}_n, \hat{\sigma}_n^2)^T = \left(\bar{x}, n^{-1} \sum_{i=1}^n (x_i - \bar{x})^2 \right)^T$



How to Learn Distributions: Toy Example

- We can explain the maximum likelihood principle using Kullback-Leibler (KL) divergence.
- Note that $n^{-1}I_n(\theta) = \int (\log p_\theta(x))p_n(x)dx$ where $p_n := n^{-1} \sum_{i=1}^n \delta_{x_i}$. It implies

$$\begin{aligned} n^{-1}I_n(\theta) &= \int (\log p_\theta(x)/p_n(x))p_n(x)dx + \int (\log p_n(x))p_n(x)dx \\ &= -\text{KL}(p_n \| p_\theta) + C \end{aligned} \tag{2}$$

where C is a constant w.r.t. θ .³

- Thus, MLEs are minimizers of the KL divergence between the empirical measure and the model density.

³Formally, $\text{KL}(p||q) := \int \log (\mathbb{P}(dx)/\mathbb{Q}(dx))\mathbb{P}(dx)$ where $\mathbb{P} \ll \mathbb{Q}$.

How to Learn Distributions: Toy Example

- There is no closed-form expression for the MLEs when addressing complex data and models.
- We usually run iterative algorithms to approximate optimal parameters.

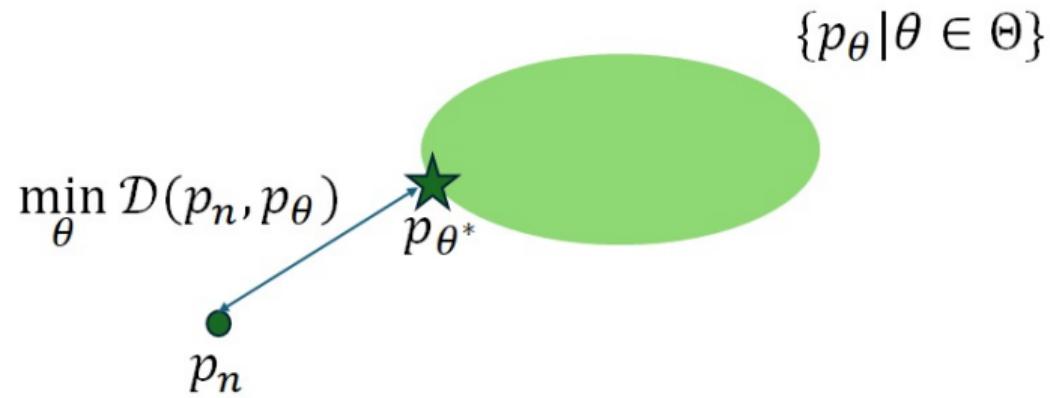
Key Elements in Generative Model

$$\{p_{\theta} | \theta \in \Theta\}$$

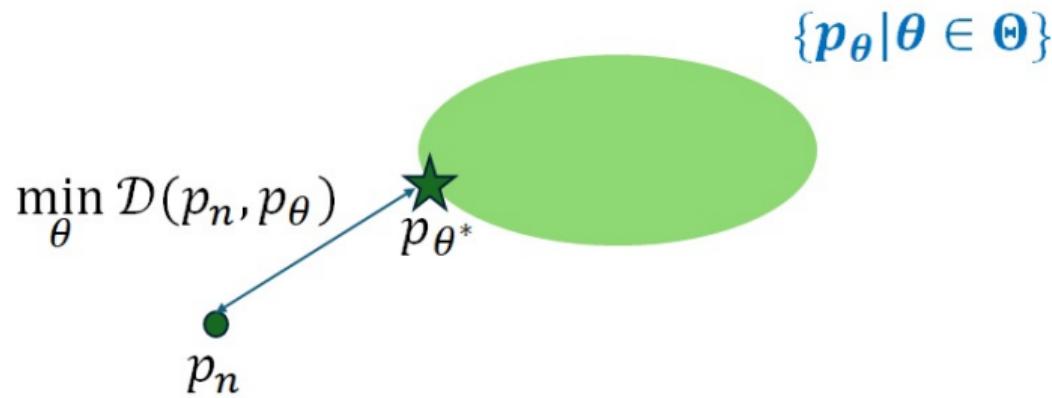


$$p_n$$

Key Elements in Generative Model



Key Elements in Generative Model



- Key elements to learn generative models include
 - ① **Model Class:** Graphical models are frequently used to reflect domain knowledge
 - ② **Statistical Distance:** Statistical distances measure differences between distributions to identify optimal generative models

Model Class



- For example, $\vec{X} := (X_1, \dots, X_m)^T \in \mathcal{X}^m$ represents a m -dimensional random vector for the 4K-resolution color images where $m \approx 24M$.⁴
- Each color channel value is discrete, ranging from 0 to 255 ($|\mathcal{X}| = 256$).
Q: How to model $p(X_1 = x_1, \dots, X_m = x_m)$?

⁴From now on, x_i represents the realization of the i -th element in \vec{X} .

Examples of Model Classes

1. **Multivariate Categorical Distribution:** Without any domain knowledge, we can introduce parameters for each realization, e.g., $p(X_1 = 0, \dots, X_m = 0)$.
 - It can express all the distributions defined on the data domain. However, the number of parameters is huge, about $|\mathcal{X}|^m$ (approximately 10^{60M}).
2. **Degenerated Model:** When (X_2, \dots, X_m) are (known) deterministic functions of X_1 , introducing parameters for the marginal distribution $p(X_1)$ is sufficient.
 - The number of parameters is small, about $|\mathcal{X}|$, and invariant to the data dimension. However, the model class is significantly reduced.

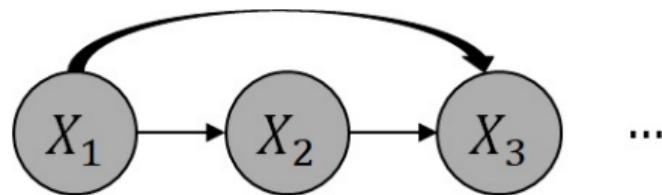
Building appropriate models that reflect domain knowledge of dependency structure is important

Model Classes by Dependency Structure

- We can express $p(X_1, \dots, X_m)$ as a product of conditional distributions:

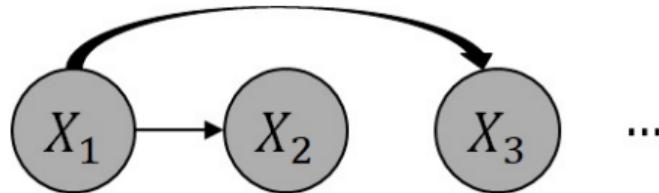
$$p(X_1)p(X_2|X_1)\dots p(X_m|X_1, X_2, \dots, X_{m-1}). \quad (3)$$

Many dependency structures can be represented by a directed graph with m nodes X_1, \dots, X_m and $(m - 1)m/2$ directed edges, such as $X_1 \rightarrow X_2, \dots, X_{m-1} \rightarrow X_m$.

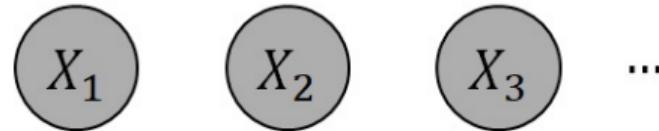


- For example, **1. Multivariate Categorical Distribution** corresponds to the graph using the all nodes and directed edges.

Model Classes by Dependency Structure



- **2. Degenerated Model:** This model is a special case of the above graph where $p_{\theta}(X_1, \dots, X_m) = p_{\theta}(X_1)p_{\theta}(X_2|X_1)\dots p_{\theta}(X_m|X_1)$.



- **3. Multivariate Independent Categorical Distribution:** This model assumes the (mutual) independency among variables, using $p_{\theta}(X_1, \dots, X_m) = p_{\theta}(X_1)\dots p_{\theta}(X_m)$. It requires $|\mathcal{X}| m$ parameters, but its assumption is strong.

Model Classes by Dependency Structure



- 4. Spatial Model:** The value of the center pixel (e.g., X_{3842}) depends only on adjacent pixels (e.g., X_1, \dots, X_{7683}), making it conditionally independent of all other pixels.
- The number of parameters is about $|\mathcal{X}|^{\# \text{ of adjacent pixels} + 1} \times m$. Assuming translation invariance reduces it to approximately $|\mathcal{X}|^{\# \text{ of adjacent pixels} + 1}$.

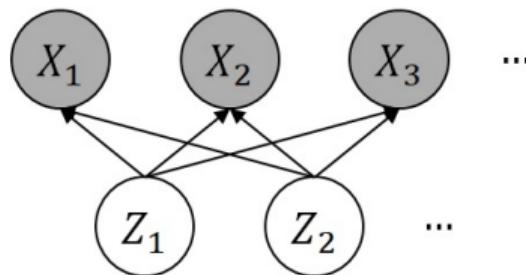
Model Classes by Dependency Structure



- 5. Latent Variable Model:** There are latent factors $\vec{Z} := (Z_1, \dots, Z_r)^T$, typically consisting of independent components, that are mixed to generate data, e.g., $\vec{X} = A\vec{Z}$ in Independent Component Analysis.
- The age, size, and location of eyes, light source location, and camera angle are examples of latent factors.

The top and bottom images are from the Extended Yale-B (Georghiades et al., 2001) and Multi-pie (Gross et al., 2010) datasets, respectively.

Model Classes by Dependency Structure

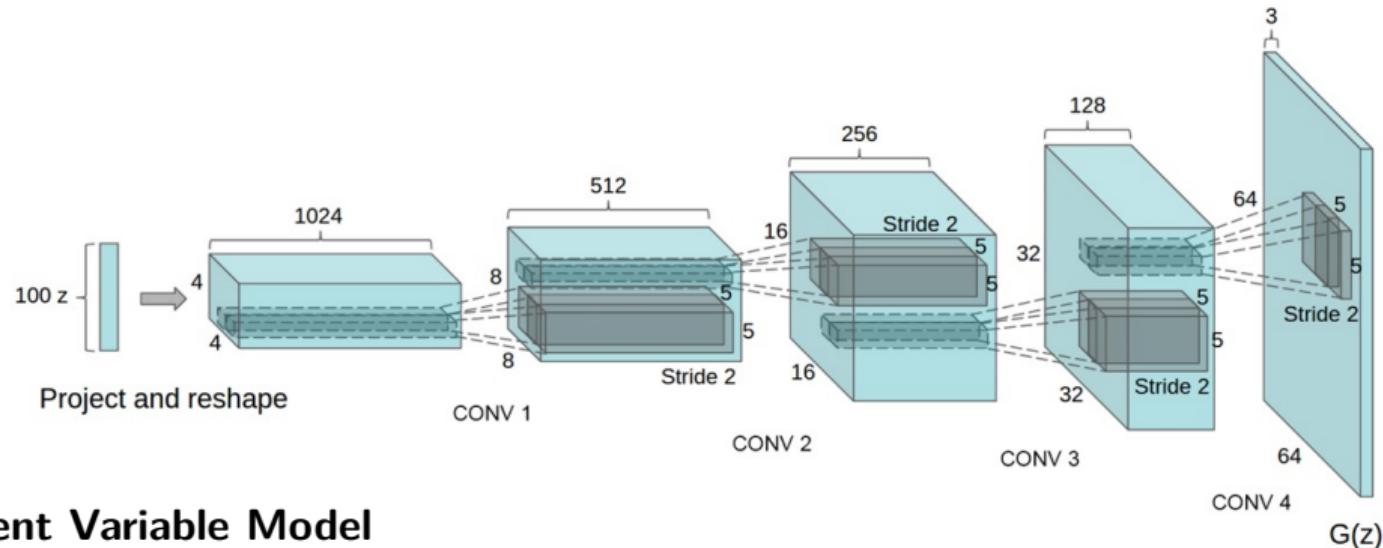


5. Latent Variable Model

$$\begin{aligned}
 p_{\theta}(X_1, \dots, X_m) &= \int \left(p_{\theta}(X_1, \dots, X_m | \vec{Z} = \vec{z}) \right) p(\vec{Z} = \vec{z}) d\vec{z} \\
 &= \int \left(p_{\theta}(X_1 | \vec{Z} = \vec{z}) \cdots p_{\theta}(X_m | \vec{Z} = \vec{z}) \right) \prod_{i=1}^r p(Z_i = z_i) d\vec{z}
 \end{aligned} \tag{4}$$

- When we have discrete \vec{Z} , the number of parameters is approximately $(|\mathcal{X}| \cdot |\mathcal{Z}|^r)m$.

Model Classes by Dependency Structure

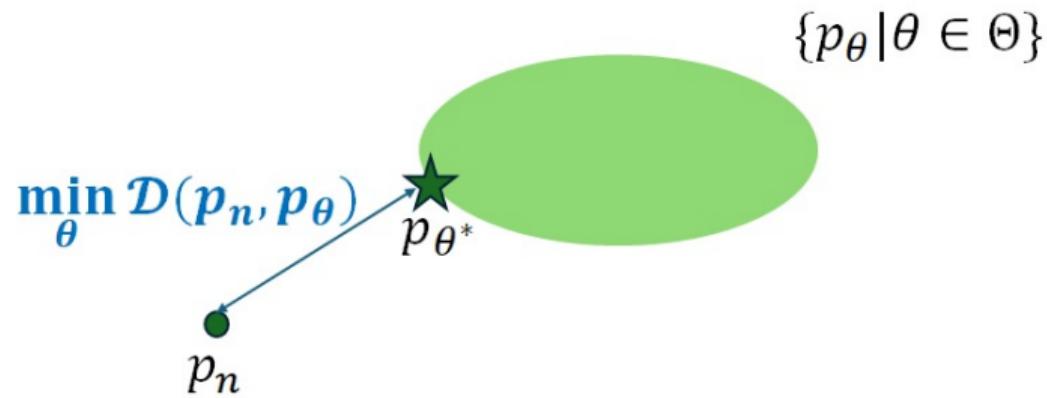


5. Latent Variable Model

- Deep generative models are predominantly based on latent variable models.
- The conditional distributions $p_{\theta}(\vec{X}|\vec{Z})$ are usually modeled as parametric family distributions, e.g., $N(\mu_{\vec{X}|\vec{Z}}(\vec{Z}), \Sigma_{\vec{X}|\vec{Z}}(\vec{Z}))$, which further reduce the number of parameters.

An example neural network for deep generative model from Radford (2015).

Key Elements in Generative Model



- Key elements to learn generative models include
 - ① **Model Class:** Graphical models are frequently used to reflect domain knowledge
 - ② **Statistical Distance:** Statistical distances measure differences between distributions to identify optimal generative models

Examples of Statistical Distances

- Deep generative models learn p_n by minimizing $\mathcal{D}(p_n, p_\theta)$, where \mathcal{D} denotes a statistical distance.
- The effectiveness of \mathcal{D} varies depending on the type of data and the specific algorithm implementation. Each \mathcal{D} necessitates different model classes and corresponding loss functions.
- Popular choices of \mathcal{D} include:
 - ① f -divergence
 - ② Integral Probability Metric
 - ③ Wasserstein Distance⁵
 - ④ Fisher Divergence

⁵Named after “Leonid Vaserštejn”, though “Wasserstein” is more commonly used in English publications.

Examples of Statistical Distances

1. *f*-divergence:

- The *f*-divergences (Rényi, 1961) are expectations of density ratios mapped by convex functions $f : \mathbb{R}^+ \rightarrow \mathbb{R}$, satisfying $f(1) = 0$. They can be expressed as:

$$\mathcal{D}_f(p||q) := \int f\left(\frac{p(\vec{x})}{q(\vec{x})}\right) q(\vec{x}) d\vec{x} \quad (5)$$

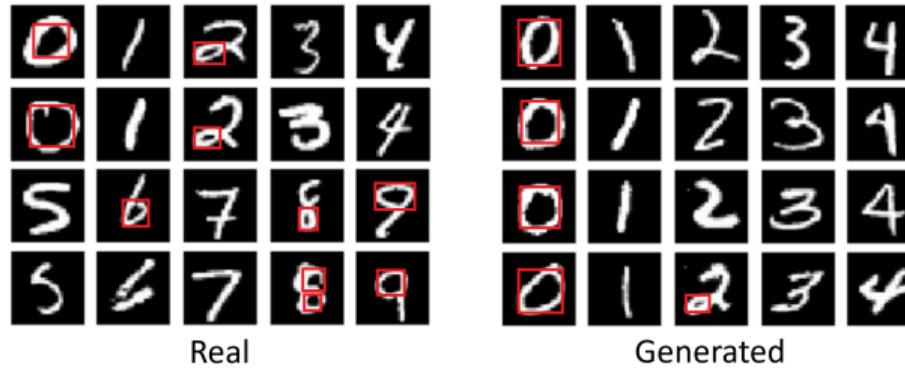
- The $\text{KL}(p_n||p_\theta)$ equals $\mathcal{D}_f(p_n||p_\theta)$ when $f(u) = u \log u$.

Proof:

$$\mathcal{D}_f(p_n||p_\theta) = \int \left(p_n(\vec{x})/p_\theta(\vec{x})\right) \log \left(p_n(\vec{x})/p_\theta(\vec{x})\right) p_\theta(\vec{x}) d\vec{x} = \int \log \left(p_n(\vec{x})/p_\theta(\vec{x})\right) p_n(\vec{x}) d\vec{x}.$$

- Thus, all maximum likelihood methods target minimizing this specific *f*-divergence.

Examples of Statistical Distances



2. **Integral Probability Metric:** The integral probability metrics (IPMs, Müller, 1997) between distributions are the largest difference between their summary statistics.
- For example, when we use the number of circles in images as summary statistics, the difference is 0.5 (real) – 0.25 (generated) = 0.25 .

Image source: MNIST (Deng, 2012)

Examples of Statistical Distances

2. Integral Probability Metric:

- We can consider many summary statistics together to precisely compare distributions.
- The IPMs can be expressed as:

$$\gamma_{\mathcal{F}}(\mathbb{P}, \mathbb{Q}) := \sup_{f \in \mathcal{F}} \left| \int f(\vec{x}) d\mathbb{P}(\vec{x}) - \int f(\vec{x}) d\mathbb{Q}(\vec{x}) \right| \quad (6)$$

where \mathcal{F} is a class of real-valued functions, and \mathbb{P} and \mathbb{Q} are probability measures.

- Examples of \mathcal{F} include the class of all 1-Lipschitz continuous functions, and functions from reproducing kernel Hilbert space (RKHS).

Examples of Statistical Distances

	0	1	2		0	1	2	
0	1/9	1/9	1/9		0	0.10	0.30	0.40
1	1/9	1/9	1/9		1	0.40	0.05	0.25
2	1/9	1/9	1/9		2	0.25	0.30	0.10
Joint Distribution				Transportation Cost				

3. **Wasserstein Distance:** Wasserstein distance (Monge, 1781; Kantorovich, 1960)

represents the minimum expected transportation cost between two distributions.

- For example, consider the above joint distribution. The transportation cost is calculated as:

$$\begin{aligned}
 & (1/3) \times (0.10 \times (1/3) + 0.30 \times (1/3) + 0.40 \times (1/3)) \\
 & + (1/3) \times (0.40 \times (1/3) + 0.05 \times (1/3) + 0.25 \times (1/3)) \\
 & + (1/3) \times (0.25 \times (1/3) + 0.30 \times (1/3) + 0.10 \times (1/3)) = 0.24.
 \end{aligned}$$

Examples of Statistical Distances

		0	1	2		0	1	2	
		0	1/3	0	0	0	0.10	0.30	0.40
		1	0	1/3	0	1	0.40	0.05	0.25
		2	0	0	1/3	2	0.25	0.30	0.10
Joint Distribution		Transportation Cost							

3. Wasserstein Distance

- We can consider another joint distribution. The transportation cost is

$$\begin{aligned}
 & (1/3)(0.10(1) + 0.30(0) + 0.40(0)) \\
 & + (1/3)(0.40(0) + 0.05(1) + 0.25(0)) \\
 & + (1/3)(0.25(0) + 0.30(0) + 0.10(1)) = 0.08.
 \end{aligned}$$

Examples of Statistical Distances

3. Wasserstein Distance

- The p -Wasserstein distance can be expressed as

$$W_p(\mathbb{P}, \mathbb{Q}; d) := \left(\inf_{\pi \in \Pi(\mathbb{P}, \mathbb{Q})} \int d^p(\vec{x}, \vec{x}') d\pi(\vec{x}, \vec{x}') \right)^{1/p} \quad (7)$$

where $p \in [1, \infty)$ and $\Pi(\mathbb{P}, \mathbb{Q})$ is the set of all joint distributions whose marginals are \mathbb{P} and \mathbb{Q} .

Examples of Statistical Distances

- 4. Fisher Divergence:** Fisher divergence (Johnson, 2004; Hyvärinen, 2005) is the expected difference between the (Stein) scores (Liu et al., 2016)⁶ of two distributions. It can be expressed as:

$$\text{FD}(p \parallel q) = \int \|\nabla_{\vec{x}} \log p(\vec{x}) - \nabla_{\vec{x}} \log q(\vec{x})\|^2 p(\vec{x}) d\vec{x}. \quad (8)$$

- It is zero if and only if $p = q$.

Proof:

$$\nabla_{\vec{x}} \log p(\vec{x}) = \nabla_{\vec{x}} \log q(\vec{x}) \implies p(\vec{x}) = Cq(\vec{x}) \quad (9)$$

and $C = 1$ because $\int p(\vec{x}) d\vec{x} = \int q(\vec{x}) d\vec{x} = 1$.

⁶The term ‘score’ here refers to the gradient w.r.t. realizations, which differs from usual terms in parametric family distributions.

Examples of Statistical Distances

4. Fisher Divergence: This is a useful measure for learning energy-based models (Teh et al., 2003), such as Boltzmann distributions:

$$p_\theta(\vec{x}) = \frac{1}{C(\theta)} \exp(-E_\theta(\vec{x})) \quad (10)$$

where $C(\theta) := \int \exp(-E_\theta(\vec{x})) d\vec{x}$. In this case,

$$\nabla_{\vec{x}} \log p_\theta(\vec{x}) = -\nabla_{\vec{x}} E_\theta(\vec{x}) \quad (11)$$

holds, and the normalizing constant disappears.

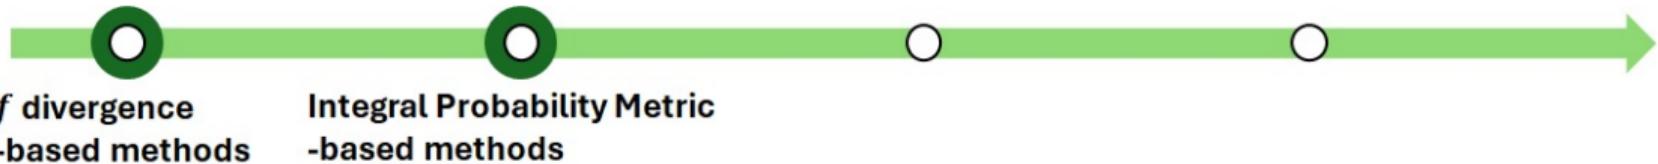
Chronicle of Deep Generative Model



f divergence
-based methods

- Variational Autoencoder (VAE, Kingma and Welling, 2014)
- Generative Adversarial Network (GAN, Goodfellow et al., 2014)
- f -GAN (Nowozin et al., 2016)

Chronicle of Deep Generative Model



- Generative Moment Matching Networks (Li et al., 2015)
- Maximum Mean Discrepancy GANs (Li et al., 2017)
- Sobolev GANs (Mroueh et al., 2017)

Chronicle of Deep Generative Model



- Wasserstein GANs (Arjovsky et al., 2017)
- Wasserstein GAN with gradient penalty (Gulrajani et al., 2017)
- Wasserstein Autoencoders (Tolstikhin et al., 2018)

Chronicle of Deep Generative Model



- Noise Conditional Score Networks (Song and Ermon, 2019)
- Denoising Diffusion Probabilistic Models (Ho et al., 2020)

Advanced Topics

- Asymptotic efficiencies according to statistical distances, e.g., minimax convergence rates of IPM-based generative models (Uppal et al., 2019).
- Methods for data domains other than images, e.g., generative pre-trained transformers (GPTs, Radford, 2018) and their variations for language data.
- Advanced graphical models reflecting domain knowledge, e.g., DALL-E (Ramesh et al., 2021) for generating images from text descriptions.

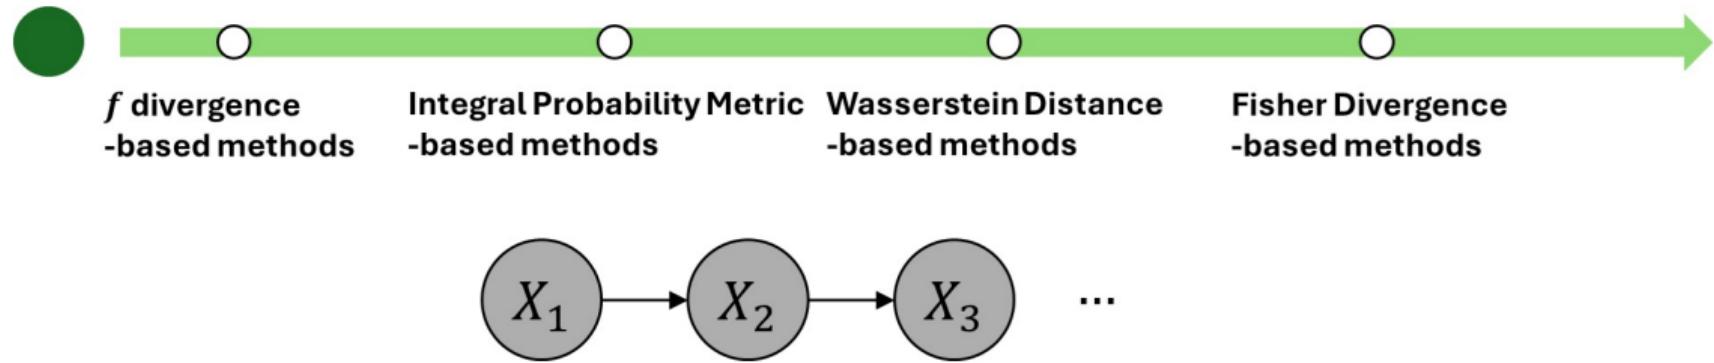
Outline

1 INTRODUCTION

2 STATISTICAL DISTANCES IN DEEP GENERATIVE MODELS

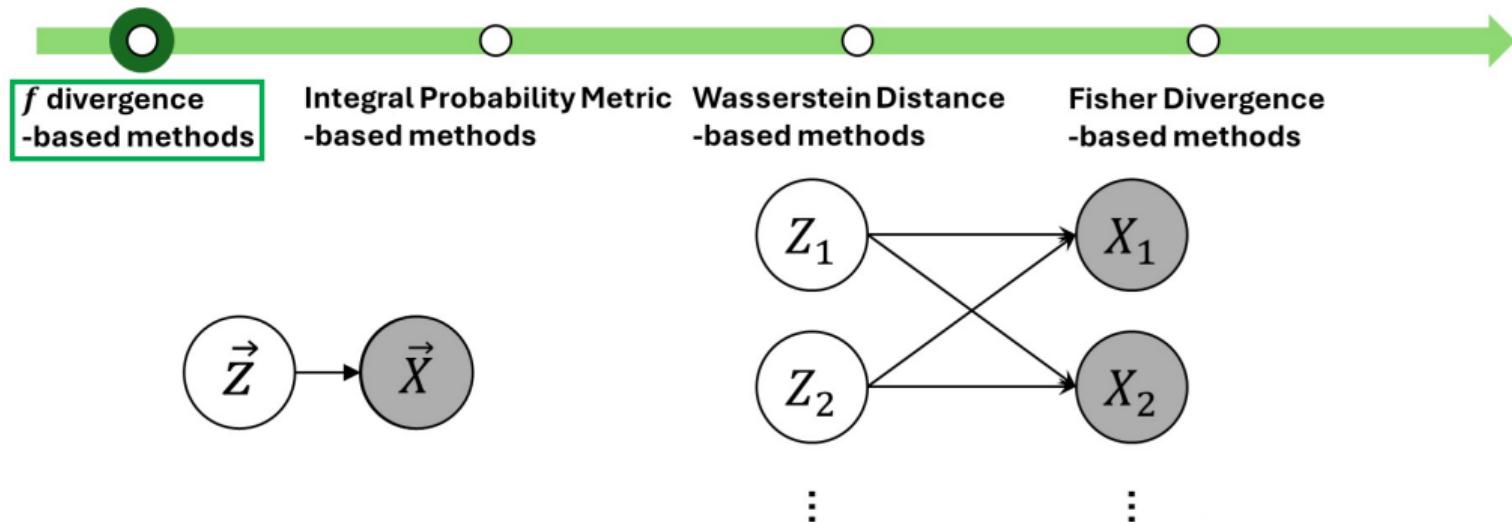
- f Divergence-based Methods
- Integral Probability Metric-based Methods
- Wasserstein Distance-based Methods
- Fisher Divergence-based Methods

Limitation of Pre-Deep Generative Model



- Before the emergence of deep generative models, state-of-the-art methods typically employed Markov models, requiring extensive Markov chain Monte Carlo computations.

Emergence of f -Divergence-based Methods



- A line of work introduced latent variable models, utilizing deep neural networks to model generator functions that mix latent variables to synthesize high-dimensional observations.
- VAEs and GANs are popular examples of these methods, specifically targeting the f -divergence, $\mathcal{D}_f(p_n \| p_\theta)$.

Recapping f -Divergence

- $\mathcal{D}_f(p||q) := \int f(p(\vec{x})/q(\vec{x})) q(\vec{x}) d\vec{x}$ where $f : \mathbb{R}^+ \rightarrow \mathbb{R}$ is a convex function satisfying $f(1) = 0$. Examples include KL divergence, total variation distance, and Jensen-Shannon (JS) divergence:
- **KL divergence:** $\text{KL}(p||q) := \int \log(p(\vec{x})/q(\vec{x})) p(\vec{x}) d\vec{x}$
- **Total variation distance:** $\delta(p, q) := \frac{1}{2} \int |p(\vec{x}) - q(\vec{x})| d\vec{x}$
- **Jensen-Shannon divergence:**

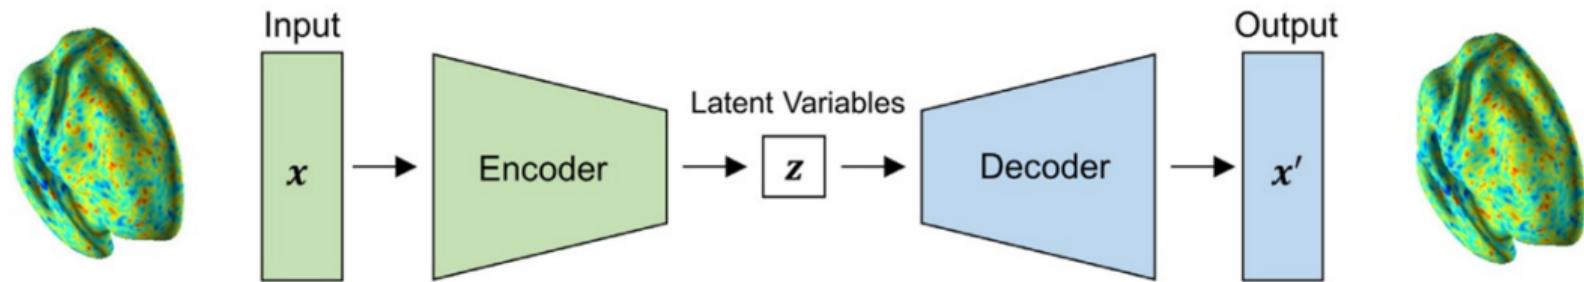
$$\text{JS}(p||q) := \frac{1}{2} \left(\text{KL}(p||\frac{p+q}{2}) + \text{KL}(q||\frac{p+q}{2}) \right) \quad (12)$$

Recapping f -Divergence

Name	$D_f(P\ Q)$	$f(u)$
Kullback-Leibler	$\int p(x) \log \frac{p(x)}{q(x)} dx$	$u \log u$
Reverse KL	$\int q(x) \log \frac{q(x)}{p(x)} dx$	$-\log u$
Pearson χ^2	$\int \frac{(q(x)-p(x))^2}{p(x)} dx$	$(u - 1)^2$
Squared Hellinger	$\int \left(\sqrt{p(x)} - \sqrt{q(x)} \right)^2 dx$	$(\sqrt{u} - 1)^2$
Jensen-Shannon	$\frac{1}{2} \int p(x) \log \frac{2p(x)}{p(x)+q(x)} + q(x) \log \frac{2q(x)}{p(x)+q(x)} dx$	$-(u+1) \log \frac{1+u}{2} + u \log u$

List of popular examples of f -divergences, edited from Nowozin et al. (2016).

KL Divergence: Variational Autoencoder

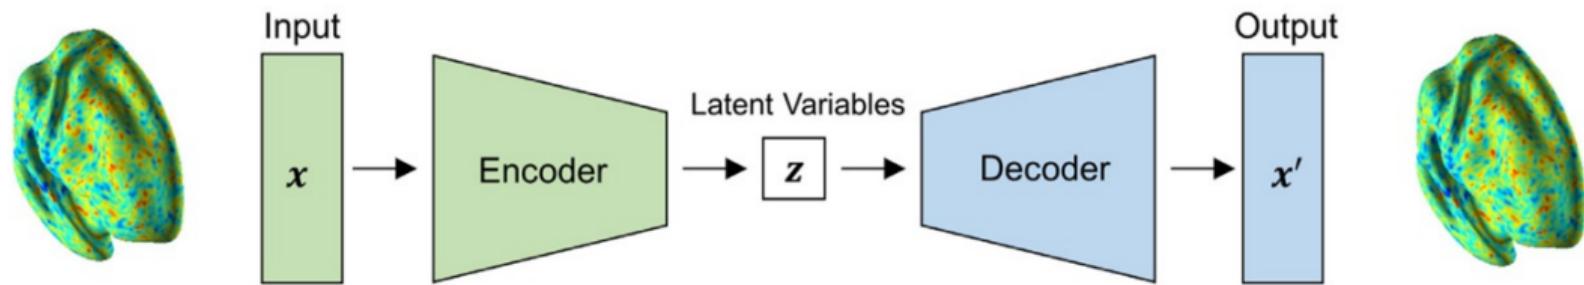


1. Variational Autoencoder:

- We first briefly review autoencoders (Bengio et al., 2006). Autoencoders (AEs) consist of pairs of encoders and decoders that efficiently reduce the dimensionality of data.
- Encoders embed observations into a lower-dimensional space (referred to as ‘encoding’), while decoders map these encodings back to the original observation space (‘decoding’ or ‘reconstruction’).

Images are from Kim et al., 2021.

KL Divergence: Variational Autoencoder



- The prefix 'auto' is used because they autonomously learn to encode data in an unsupervised manner.
- Autoencoders (AEs) are trained by minimizing the difference between the original observations and their reconstructions, referred to as the 'reconstruction error'.

Images are from Kim et al., 2021.

KL Divergence: Variational Autoencoder

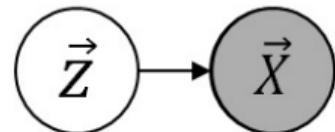
- AEs are nonlinear extensions of Principal Component Analysis (Kramer, 1991; Plaut, 2018).
- Assuming the data $(\vec{x}_i)_{i=1}^n$ is centered, for a given dimension r , we define:

$$W^* \in \arg \min_W \left(n^{-1} \sum_{i=1}^n \|\vec{x}_i - WW^T \vec{x}_i\|^2 \right) \text{ subject to } W^T W = I_r. \quad (13)$$

Here, $W^T \vec{x}_i$ represents the encoding process, and $WW^T \vec{x}_i$ represents the decoding.

- The W^* identifies optimal linear encoder and decoder pairs among symmetric AEs.
- The optimal embeddings $W^{*T} \vec{x}_i$ are the first r principal components up to orthogonal transformations.

KL Divergence: Variational Autoencoder



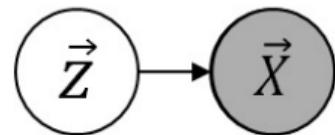
- Variational Autoencoders (VAEs, Kingma and Welling, 2014) model the data generation process using decoder networks:

$$\begin{aligned}
 p_{\theta}(\vec{z}, \vec{x}) &:= p(\vec{z})p_{\theta}(\vec{x}|\vec{z}) \\
 &= \left(\prod_{i=1}^r p(z_i) \right) \left(\prod_{i=1}^m p_{\theta}(x_i|\vec{z}) \right).
 \end{aligned} \tag{14}$$

where $p(\vec{z})$ is called the ‘prior’ distribution.

- Let $p_{\theta}(\vec{x}) := \int p_{\theta}(\vec{z}, \vec{x})d\vec{z}$. Then, $p_{\theta}(\vec{z}|\vec{x})$ represents the ‘posterior’ distribution.

KL Divergence: Variational Autoencoder



- VAEs use $N(0, I)$ for $p(\vec{z})$ and $N(\mu_{\vec{X}|\vec{Z}}(\vec{Z}), \mathbf{D}_{\vec{X}|\vec{Z}}(\vec{Z}))$ for $p_\theta(\vec{x}|\vec{z})$. Here, $\mu_{\vec{X}|\vec{Z}}$ is defined as $(\mu_{X_1|\vec{Z}}, \dots, \mu_{X_m|\vec{Z}})^T$ and $\mathbf{D}_{\vec{X}|\vec{Z}}$ as $\text{diag}(\sigma_{X_1|\vec{Z}}^2, \dots, \sigma_{X_m|\vec{Z}}^2)$, and all elements are outputs of neural networks parameterized by θ .
- In this context, the joint distribution is given by:

$$p_\theta(\vec{z}, \vec{x}) = \prod_{i=1}^r \left(\frac{1}{\sqrt{2\pi}} \exp\left(-\frac{z_i^2}{2}\right) \right) \prod_{i=1}^m \left(\frac{1}{\sqrt{2\pi\sigma_{X_i|\vec{Z}}^2(\vec{z})}} \exp\left(-\frac{(x_i - \mu_{X_i|\vec{Z}}(\vec{z}))^2}{2\sigma_{X_i|\vec{Z}}^2(\vec{z})}\right) \right). \quad (15)$$

KL Divergence: Variational Autoencoder

- Since \vec{Z} is unobserved, VAEs target to maximize $p_\theta(\vec{x}) := \int p(\vec{z})p_\theta(\vec{x}|\vec{z})d\vec{z}$. However, the $p_\theta(\vec{x})$ does not have a closed-form expression.
- VAEs apply variational inference (Bishop, 2006), introducing encoder to approximate the posterior as $q_\phi(\vec{z}|\vec{x})$. They maximize evidence lower bound (ELBO), a lower bound of the evidence $\log p_\theta(\vec{x})$:

$$\begin{aligned} \text{ELBO}(\theta, \phi; \vec{x}) &:= \log p_\theta(\vec{x}) - \text{KL}(q_\phi(\vec{z}|\vec{x}) || p_\theta(\vec{z}|\vec{x})) \\ &= \int (\log p_\theta(\vec{x}|\vec{z})) q_\phi(\vec{z}|\vec{x}) d\vec{z} - \text{KL}(q_\phi(\vec{z}|\vec{x}) || p(\vec{z})). \end{aligned} \tag{16}$$

Proof: By Bayes' rule, the relation $p_\theta(\vec{x}) = p_\theta(\vec{x}|\vec{z})p(\vec{z})/p_\theta(\vec{z}|\vec{x})$ holds, implying $\log p_\theta(\vec{x}) - \log (q_\phi(\vec{z}|\vec{x})/p_\theta(\vec{z}|\vec{x})) = \log p_\theta(\vec{x}|\vec{z}) - \log (q_\phi(\vec{z}|\vec{x})/p(\vec{z}))$. Taking the expectation over $q_\phi(\vec{z}|\vec{x})$ concludes the proof.

KL Divergence: Variational Autoencoder

- VAEs maximize the average of the ELBO, which is equivalent to minimizing

$$-\int \text{ELBO}(\theta, \phi; \vec{x}) p_n(\vec{x}) d\vec{x}. \quad (17)$$

- Define $\theta^* \in \arg \min_{\theta} \left(\min_{\phi} \left(- \int \text{ELBO}(\theta, \phi; \vec{x}) p_n(\vec{x}) d\vec{x} \right) \right)$. Then, p_{θ^*} is a minimizer of $\text{KL}(p_n \| p_{\theta^*})$.
- We assume that the encoder class $\{q_{\phi} | \phi \in \Phi\}$ is sufficiently flexible such that for any given θ , there exists a $\phi^*(\theta)$ where: $q_{\phi^*(\theta)}(\vec{z}|\vec{x}) = p_{\theta}(\vec{z}|\vec{x})$ (a.s. w.r.t. $p_n(\vec{x})$).

KL Divergence: Variational Autoencoder

Proof: By Equation (16),

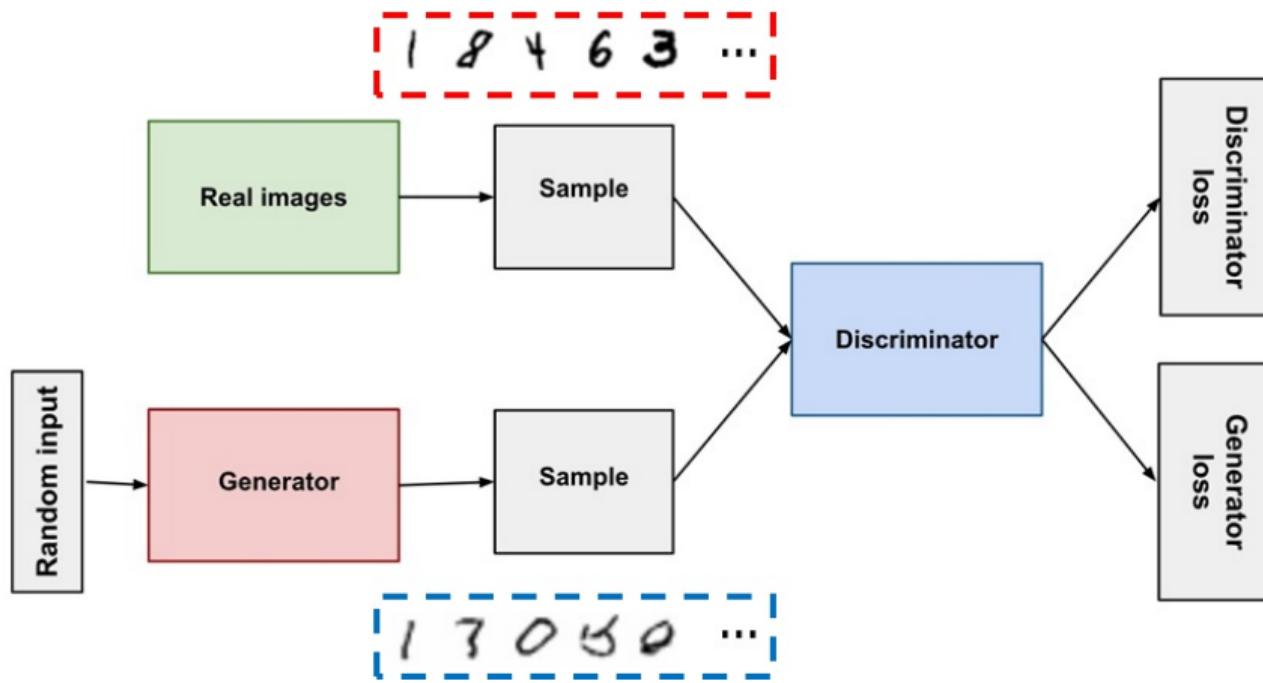
$$\begin{aligned}
 & \min_{\phi} \left(- \int \text{ELBO}(\theta, \phi; \vec{x}) p_n(\vec{x}) d\vec{x} \right) \\
 &= \min_{\phi} \left(- \int \left(\log p_\theta(\vec{x}) - \text{KL}(q_\phi(\vec{z}|\vec{x}) || p_\theta(\vec{z}|\vec{x})) \right) p_n(\vec{x}) d\vec{x} \right) \\
 &= \text{KL}(p_n || p_\theta) + \min_{\phi} \int \text{KL}(q_\phi(\vec{z}|\vec{x}) || p_\theta(\vec{z}|\vec{x})) p_n(\vec{x}) d\vec{x} + C.
 \end{aligned}$$

Thus, $\min_{\phi} \left(- \int \text{ELBO}(\theta, \phi; \vec{x}) p_n(\vec{x}) d\vec{x} \right) = - \int \text{ELBO}(\theta, \phi^*(\theta); \vec{x}) p_n(\vec{x}) d\vec{x} = \text{KL}(p_n || p_\theta)$ up to a constant addition.

JS Divergence: Generative Adversarial Network

2. **Generative Adversarial Network:** Generative Adversarial Networks (GANs; Goodfellow et al., 2014) introduce adversarial learning through two networks: a generator and a discriminator.
- The generator specifies the same graphical model used in VAEs, but $\vec{x} = G_{\theta}(\vec{z})$ where G is a neural network, i.e., $p_{\theta}(\vec{x}|\vec{z})$ degenerates to a single point.
 - The discriminator is a binary classifier designed to differentiate between real data and synthetic data produced by the generator.

JS Divergence: Generative Adversarial Network



Images were edited from https://developers.google.com/machine-learning/gan/gan_structure and <https://github.com/MorvanZhou/mnistGANs>.

JS Divergence: Generative Adversarial Network

- The adversarial learning process involves alternately maximizing and minimizing the negative cross-entropy loss:

$$V(\theta, \phi) := \int (\log D_\phi(\vec{x})) p_n(\vec{x}) d\vec{x} + \int (\log(1 - D_\phi(\vec{x}))) p_\theta(\vec{x}) d\vec{x}. \quad (18)$$

- This process can be viewed as a two-player minimax game where the goal is to find

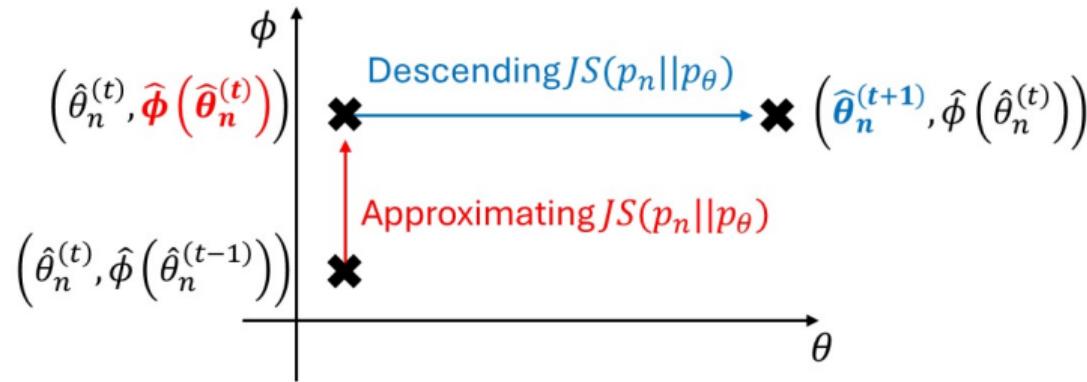
$$\theta^* \in \arg \min_{\theta} \left(\max_{\phi} V(\theta, \phi) \right). \quad (19)$$

- The adversarial training consists of repeated cycles of approximating and minimizing the JS divergence:

$$\max_{\phi} V(\theta, \phi) = \text{JS}(p_n \parallel p_\theta) \quad (20)$$

up to a constant addition and sign-preserving multiplication. Thus, p_{θ^*} is the minimizer of $\text{JS}(p_n \parallel p_\theta)$.

JS Divergence: Generative Adversarial Network



- The two adversarial networks, the discriminator and the generator, are trained alternately:
 - Given $\hat{\theta}_n^{(t)}$, update the discriminator to obtain $\hat{\phi}(\hat{\theta}_n^{(t)})$ by maximizing $V(\hat{\theta}_n^{(t)}, \phi)$.
 - Given $\hat{\phi}(\hat{\theta}_n^{(t)})$, update the generator to obtain $\hat{\theta}_n^{(t+1)}$ by minimizing $V(\theta, \hat{\phi}(\hat{\theta}_n^{(t)}))$.
 - Repeat the above processes.

JS Divergence: Generative Adversarial Network

- Again, $\max_{\phi} V(\theta, \phi) = \text{JS}(p_n \| p_\theta)$ up to trivial transformations. We assume that the discriminator class $\{D_\phi \mid \phi \in \Phi\}$ is sufficiently flexible such that for any given θ , there exists a $\phi^*(\theta)$ where:

$$D_{\phi^*(\theta)}(\vec{x}) = \frac{p_n(\vec{x})}{p_n(\vec{x}) + p_\theta(\vec{x})}. \quad (21)$$

Proof: $V(\theta, \phi) = \int \left(\log D_\phi(\vec{x}) p_n(\vec{x}) + \log(1 - D_\phi(\vec{x})) p_\theta(\vec{x}) \right) d\vec{x}$, and the integrand is strictly concave w.r.t. $D_\phi(\vec{x})$. The first derivative of the integrand w.r.t. $D_\phi(\vec{x})$ is

$$-\frac{p_n(\vec{x}) + p_\theta(\vec{x})}{D_\phi(\vec{x})(1 - D_\phi(\vec{x}))} \left(D_\phi(\vec{x}) - \frac{p_n(\vec{x})}{p_n(\vec{x}) + p_\theta(\vec{x})} \right), \quad (22)$$

implying that $V(\theta, \phi)$ is maximized when Equation (21) holds. This implies

$$\max_{\phi} V(\theta, \phi) = V(\theta, \phi^*(\theta)) = 2\text{JS}(p_n \| p_\theta) - \log 4. \quad (23)$$

f -Divergence: f -GAN

3. f -GAN:

- We have reviewed the following relationship in GANs, which holds up to a constant addition and sign-preserving multiplication: Using discriminator networks parameterized with ϕ ,

$$\text{JS}(p_n \parallel p_\theta) \approx V(\theta, \hat{\phi}_n(\theta)).$$

- Nowozin et al. (2016) generalized the concept of using auxiliary networks to approximate other f -divergences.
- The key idea is to introduce the convex conjugate function (or Fenchel conjugate, Hiriart-Urruty and Lemaréchal, 2004) to derive variational estimations of f -divergences.

f -Divergence: f -GAN

- We denote the convex conjugates of functions f by $f^*(t) := \sup_u\{ut - f(u)\}.$
- The f^* relates f and its subgradients. When f is convex and differentiable⁷, the following properties hold:
 - ① f^* is also convex and differentiable.
 - ② Duality holds, i.e., $(f^*)^* = f$.
 - ③ The relation $f(u) + f^*(t) = ut$ holds if and only if $t = f'(u)$.
- When f' is invertible, $f^*(t) = (f')^{-1}(t)t - (f \circ f'^{-1})(t)$.

⁷For more general functions, check Hiriart-Urruty and Lemaréchal (2004).

f -Divergence: f -GAN

- By duality and the definition of supremum, we have the following variational formulation:

$$\begin{aligned} \mathcal{D}_f(p_n || p_\theta) &= \int \sup_t \left\{ \frac{p_n(\vec{x})}{p_\theta(\vec{x})} t - f^*(t) \right\} p_\theta(\vec{x}) d\vec{x} \\ &\geq \sup_{T_\phi} \left(\int T_\phi(\vec{x}) p_n(\vec{x}) d\vec{x} - \int f^*(T_\phi(\vec{x})) p_\theta(\vec{x}) d\vec{x} \right). \end{aligned} \quad (24)$$

- We assume that f is differentiable and that $\{T_\phi | \phi \in \Phi\}$ is sufficiently flexible such that for any given θ , there exists $\phi^*(\theta)$ where:

$$T_{\phi^*(\theta)}(\vec{x}) = f'(p_n(\vec{x})/p_\theta(\vec{x})). \quad (25)$$

This satisfies the equality condition of Equation (24).

f -Divergence: f -GAN

- Define $F(\theta, \phi) := \int T_\phi(\vec{x}) p_n(\vec{x}) d\vec{x} - \int f^*(T_\phi(\vec{x})) p_\theta(\vec{x}) d\vec{x}$ and $\theta^* := \arg \min_{\theta} \left(\max_{\phi} F(\theta, \phi) \right)$.
- Then, $\max_{\phi} F(\theta, \phi) = F(\theta, \phi^*(\theta)) = \mathcal{D}_f(p_n \| p_\theta)$. Thus, p_{θ^*} is a minimizer of the f -divergence.
- **Example 1 (KL divergence):** Let $f(u) = u \log u$ and $f^*(t) = \exp(t - 1)$. We can express $F(\theta, \phi)$ as follows:

$$F(\theta, \phi) = \int T_\phi(\vec{x}) p_n(\vec{x}) d\vec{x} - \int \exp(T_\phi(\vec{x}) - 1) p_\theta(\vec{x}) d\vec{x}. \quad (26)$$

f -Divergence: f -GAN

- **Example 2 (JS divergence):** Let $f(u) = -(u+1)\log\frac{1+u}{2} + u\log u$ and $f^*(t) = -\log(2 - \exp(t))$. We can express $F(\theta, \phi)$ as follows:

$$F(\theta, \phi) = \int T_\phi(\vec{x}) p_n(\vec{x}) d\vec{x} + \int \log\left(2 - \exp\left(T_\phi(\vec{x})\right)\right) p_\theta(\vec{x}) d\vec{x}. \quad (27)$$

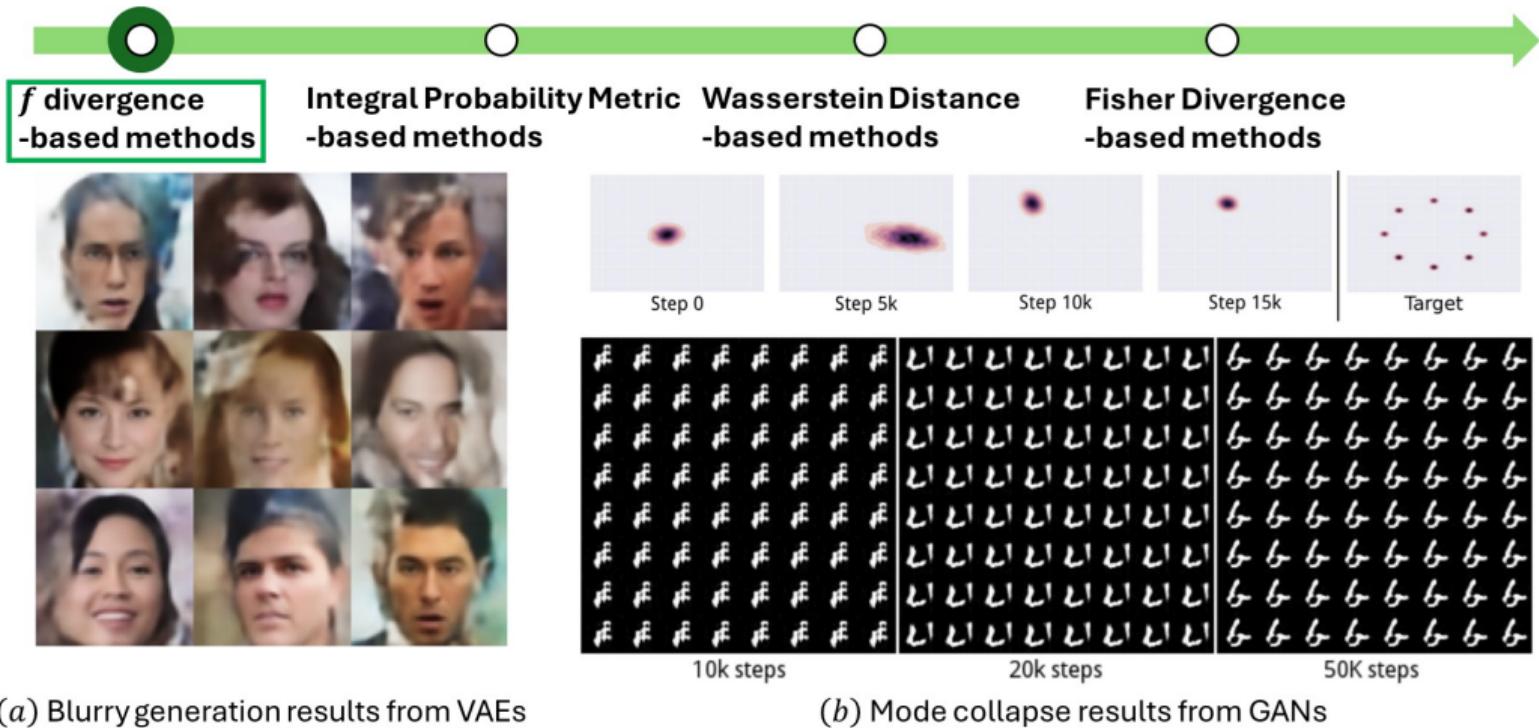
- GANs are special cases of f -GANs. When we model $T_\phi(\vec{x}) = \log D_\phi(\vec{x}) + \log 2$, $F(\theta, \phi) = V(\theta, \phi) + \log 4$, and $T_{\phi^*(\theta)}(\vec{x}) = \log D_{\phi^*(\theta)}(\vec{x}) + \log 2 = \log \frac{p_n(\vec{x})}{p_n(\vec{x}) + p_\theta(\vec{x})} + \log 2$ hold.

Generation Results: Latent Manifold Learned by VAEs



Images are edited from Kingma and Welling (2014).

Limitation of f Divergence-based Methods



Images are edited from Tolstikhin et al. (2018) and Metz et al. (2017).

Limitations of f -Divergence-Based Methods



- The mode collapse phenomenon in GANs demonstrates that the p_θ fails in capturing the support of p_n .
- Several works have criticized f -divergence,

$$\mathcal{D}_f(p_n \| p_\theta) = \int f \left(\frac{p_n(\vec{x})}{p_\theta(\vec{x})} \right) p_\theta(\vec{x}) d\vec{x},$$

pointing out that it is based on the density ratio p_n/p_θ , and this dependency may be a reason for the observed failures in f -divergence-based methods.

Emergence of IPM and Wasserstein Distances-Based Methods



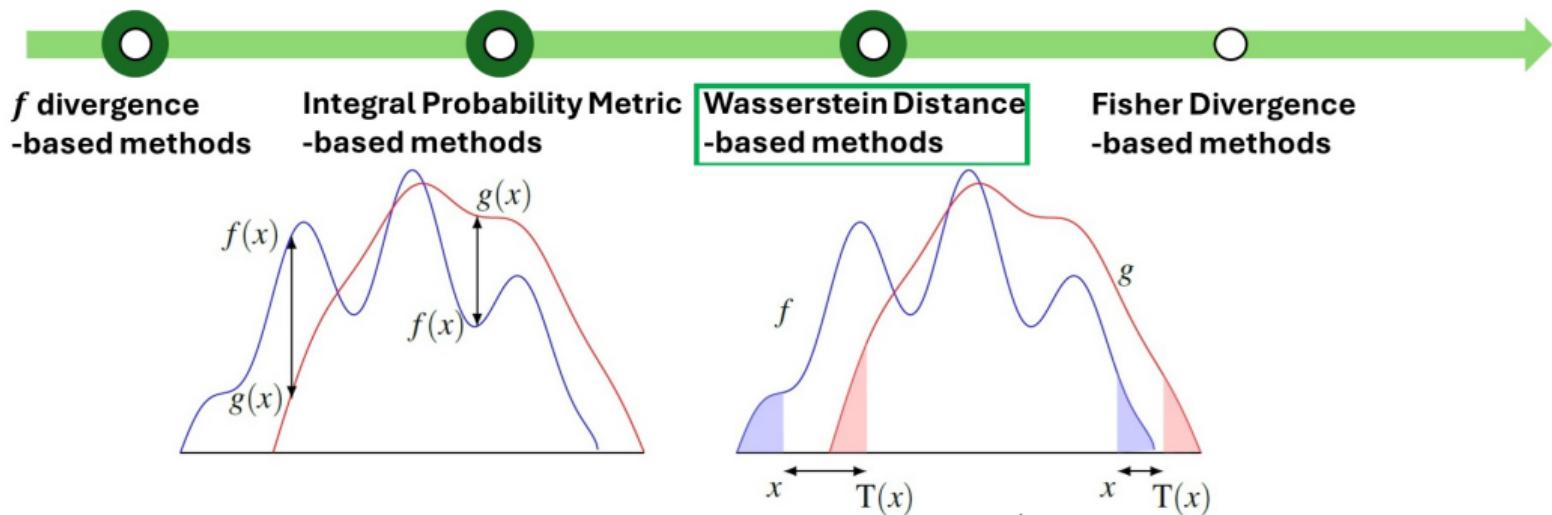
- As an alternative to density ratios, a line of work has proposed focusing on discrepancy measures that are effective regardless of the differences between the supports of p_n and p_θ .
- For example, Generative Moment Matching Networks (Li et al., 2015) aim to minimize:

$$\left\| \int \varphi(\vec{x}) d\mathbb{P}_n(\vec{x}) - \int \varphi(\vec{x}) d\mathbb{P}_\theta(\vec{x}) \right\|^2 \quad (28)$$

where the integrated terms $\varphi(\vec{x})$ represent vectors of finite moments, e.g., $\varphi(x) = (c, \sqrt{2}cx, x^2)^T$ in the univariate case with second-order moments.

- This loss function is a special case of integral probability metrics, $\gamma_{\mathcal{F}}(p_n, p_\theta)$, where \mathcal{F} denotes a set of summary statistics functions, such as moments.

Emergence of IPM and Wasserstein Distances-based Methods



- Another line of work has targeted $\left(\int d^p(\vec{x}, T(\vec{x})) d\vec{x} \right)^{1/p}$ where T transports data points from the initial distribution to the target distribution.
- This concept can be formulated as minimizing Wasserstein distances $W_p(p_n, p_\theta)$.

Images are edited from Santambrogio (2015).

Recapping Integral Probability Metric

- The IPMs can be expressed as:

$$\gamma_{\mathcal{F}}(\mathbb{P}, \mathbb{Q}) := \sup_{f \in \mathcal{F}} \left| \int f(\vec{x}) d\mathbb{P}(\vec{x}) - \int f(\vec{x}) d\mathbb{Q}(\vec{x}) \right|$$

where \mathcal{F} is a class of real-valued functions, and \mathbb{P} and \mathbb{Q} are probability measures.

Recapping Integral Probability Metric

- **Total Variation Distance:** The total variation distance, $\delta(p, q) = \frac{1}{2} \int |p(\vec{x}) - q(\vec{x})| d\vec{x}$, has an alternative expression:

$$\sup_{A \in \mathcal{A}} |\mathbb{P}(A) - \mathbb{Q}(A)| = \sup_{A \in \mathcal{A}} \left| \int I(\vec{x} \in A) d\mathbb{P}(\vec{x}) - \int I(\vec{x} \in A) d\mathbb{Q}(\vec{x}) \right|$$

where \mathcal{A} is the corresponding σ -algebra. Thus, the total variation is the IPM using the set of indicator functions for all events.

- **Earth Mover's Distance:** When \mathcal{F} consists of all 1-Lipschitz continuous functions, $\gamma_{\mathcal{F}}(\mathbb{P}, \mathbb{Q})$ corresponds to the Earth mover's distance (or 1-Wasserstein distance), a special case of Wasserstein distances. Further details will be discussed in the subsequent subsection on Wasserstein distances.

Recapping Integral Probability Metric

- **Maximum Mean Discrepancy (MMD):** We denote the kernel mean by $\mu_{\mathbb{P}}(\vec{x}) := \int k(\vec{x}', \vec{x}) d\mathbb{P}(\vec{x}')$. Then, the MMD is defined as the difference between kernel means in \mathcal{H} , the RKHS specified by k :

$$\text{MMD}_k(\mathbb{P}, \mathbb{Q}) := \|\mu_{\mathbb{P}} - \mu_{\mathbb{Q}}\|_{\mathcal{H}}.$$

MMD builds a kernel-based test statistic for a two-sample test:

$$H_0 : \mathbb{P} = \mathbb{Q} \text{ vs. } H_1 : \mathbb{P} \neq \mathbb{Q}.$$

- The MMD has important alternative representations:

$$\textcircled{1} \textbf{ IPM: } \text{MMD}_k(\mathbb{P}, \mathbb{Q}) = \sup_{\|f\|_{\mathcal{H}} \leq 1} \left(\int f(\vec{x}) d\mathbb{P}(\vec{x}) - \int f(\vec{x}) d\mathbb{Q}(\vec{x}) \right).$$

$$\textcircled{2} \textbf{ Kernel function form: }$$

$$\begin{aligned} & \text{MMD}_k^2(\mathbb{P}, \mathbb{Q}) \\ &= \int k(\vec{x}, \vec{x}') d\mathbb{P}(\vec{x}) d\mathbb{P}(\vec{x}') - 2 \int k(\vec{x}, \vec{y}) d\mathbb{P}(\vec{x}) d\mathbb{Q}(\vec{y}) + \int k(\vec{y}, \vec{y}') d\mathbb{Q}(\vec{y}) d\mathbb{Q}(\vec{y}'). \end{aligned} \tag{29}$$

MMD: Generative Moment Matching Network

1. **Generative Moment Matching Network (GMMN)**: GMMNs (Li et al., 2015) propose to use empirical estimators as loss functions to train generative models rather than introducing adversarial networks as in GANs.
- Given $(\vec{x}_i)_{i=1}^B$ and $(G_\theta(\vec{z}_i))_{i=1}^B$, minibatch samples of size B from \mathbb{P}_n and \mathbb{P}_θ respectively, the minibatch-based empirical estimators for $\text{MMD}_k^2(\mathbb{P}_n, \mathbb{P}_\theta)$ can be expressed as

$$\begin{aligned} & \frac{1}{B(B-1)} \sum_{i=1}^B \sum_{j \neq i}^B k(\vec{x}_i, \vec{x}_j) - \frac{2}{B^2} \sum_{i=1}^B \sum_{j=1}^B k(\vec{x}_i, G_\theta(\vec{z}_j)) \\ & + \frac{1}{B(B-1)} \sum_{i=1}^B \sum_{j \neq i}^B k(G_\theta(\vec{z}_i), G_\theta(\vec{z}_j)). \end{aligned} \tag{30}$$

- GMMNs used a mixture of multiple Gaussian kernels with various bandwidth parameters.

MMD: Generative Moment Matching Network

- Minimizing $\text{MMD}_k(\mathbb{P}_n, \mathbb{P}_\theta)$ can be interpreted as matching moments between \mathbb{P}_n and \mathbb{P}_θ .
- Let k be the kernel that defines the MMD, and let $\varphi(\vec{x})^8$ represent the corresponding kernel feature mapping, i.e.,

$$k(\vec{x}, \vec{x}') = \varphi(\vec{x})^\top \varphi(\vec{x}') \quad (31)$$

- For a univariate example, consider $k(x, x') = (xx' + c)^2$ for some $c > 0$. The feature mapping $\varphi(x) = (c, \sqrt{2cx}, x^2)^\top$ satisfies Equation (31). Kernels with higher degrees allow for covering higher-order moments.
- The loss of GMMNs, minibatch-based empirical estimators for (squared) MMD, can be expressed as

$$\|B^{-1} \sum_{i=1}^B \varphi(\vec{x}_i) - B^{-1} \sum_{i=1}^B \varphi(G_\theta(\vec{z}_i))\|^2. \quad (32)$$

⁸The symbol ϕ is more commonly used, but we use φ here to avoid confusion with parameters for auxiliary networks, e.g., the discriminator in GANs.

MMD: MMD GAN

2. MMD GAN:

- GMMNs face challenges in selecting effective kernels. MMD GANs (Li et al., 2017) overcome this limitation by introducing adversarial kernel learning.
- MMD GANs aim to target $\max_{k \in \mathcal{K}} \text{MMD}_k(\mathbb{P}_n, \mathbb{P}_\theta)$, where \mathcal{K} is a class of kernel functions.
- To model an expressive class \mathcal{K} , MMD GANs employ a neural network E_ϕ to define $(k \circ E_\phi)(\vec{x}, \vec{x}') := k(E_\phi(\vec{x}), E_\phi(\vec{x}'))$, targeting:

$$\max_{\phi} \text{MMD}_{k \circ E_\phi}(\mathbb{P}_n, \mathbb{P}_\theta). \quad (33)$$

- The injectivity of E_ϕ is crucial to retain the important properties of MMDs with usual kernels. MMD-GANs incorporate an encoder architecture for E_ϕ , add a decoder, and introduce a reconstruction error-based penalty term to enforce the injectivity.

Other IPMs

3. Methods using Other IPMs:

- One of the main challenges in using IPMs,

$$\gamma_{\mathcal{F}}(\mathbb{P}_n, \mathbb{P}_{\theta}) := \sup_{f \in \mathcal{F}} \left| \int f(\vec{x}) d\mathbb{P}_n(\vec{x}) - \int f(\vec{x}) d\mathbb{P}_{\theta}(\vec{x}) \right|,$$

lies in approximating the supremum over the function class \mathcal{F} .

- While MMD has a tractable representation that allows for the direct use of its empirical estimators, this is not the case for more general IPMs.
- Most methods targeting other IPMs employ neural networks to model elements within \mathcal{F} . Notably, Wasserstein GANs (Arjovsky et al., 2017) have become one of the most popular methods targeting the 1-Wasserstein distance.

Recapping Wasserstein Distance

- The p -Wasserstein distance can be expressed as

$$W_p(\mathbb{P}, \mathbb{Q}; d) := \left(\inf_{\pi \in \Pi(\mathbb{P}, \mathbb{Q})} \int d^p(\vec{x}, \vec{x}') d\pi(\vec{x}, \vec{x}') \right)^{1/p}$$

where $p \in [1, \infty)$, and $\Pi(\mathbb{P}, \mathbb{Q})$ is the set of all joint distributions whose marginals are \mathbb{P} and \mathbb{Q} .

- (Monge-Kantorovich transportation problem) Under some conditions, there exists a map T that satisfies

- 1 $W_p(\mathbb{P}, \mathbb{Q}; d) = \left(\int d^p(\vec{x}, T(\vec{x})) d\mathbb{P}(\vec{x}) \right)^{1/p}$
- 2 $\mathbb{P}(T(\vec{x})) = \mathbb{Q}(\vec{x})$ ⁹

The map T is called the ‘optimal transport map’.

⁹This can be expressed with the push-forward operation $T \# \mathbb{P} = \mathbb{Q}$.

Recapping Wasserstein Distance

- For example, when d is the Euclidean norm, W_p becomes the Mallows metric (Mallows, 1972), and has played an important role in deriving asymptotic properties of bootstrap estimators (Bickel and Freedman, 1981; Freedman, 1981).
- When $p = 1$, duality holds (Villani et al., 2009; Villani, 2021), which provides an IPM formulation:

$$W_1(\mathbb{P}, \mathbb{Q}; d) = \sup_{\|f\|_L \leq 1} \int f(\vec{x}) d\mathbb{P}(\vec{x}) - \int f(\vec{x}) d\mathbb{Q}(\vec{x}) \quad (34)$$

where $\|f\|_L := \max\{C |f(\vec{x}) - f(\vec{x}')| \leq Cd(\vec{x}, \vec{x}')\}$ represents the Lipschitz constant of f .

Recapping Wasserstein Distance

- Wasserstein distances effectively quantify differences between high-dimensional distributions when their supports are in low-dimensional manifolds.

Example

(Example 1 in Arjovsky et al., 2017) Let $Z \sim U[0, 1]$, $X = (0, Z)$, and $G_\theta(Z) = (\theta, Z)$.

- Intuitively, $\mathcal{D}(\mathbb{P}_{n=\infty}, \mathbb{P}_\theta)$ should decrease as θ vanishes.
 - $W_p(\mathbb{P}_{n=\infty}, \mathbb{P}_\theta; |\cdot|) = |\theta|$
 - $\text{JS}(\mathbb{P}_{n=\infty} \parallel \mathbb{P}_\theta) = \log 2$ if $\theta \neq 0$ and 0 if $\theta = 0$
 - $\text{KL}(\mathbb{P}_{n=\infty} \parallel \mathbb{P}_\theta) = \infty$ if $\theta \neq 0$ and 0 if $\theta = 0$
 - $\delta(\mathbb{P}_{n=\infty}, \mathbb{P}_\theta) = 1$ if $\theta \neq 0$ and 0 if $\theta = 0$

1-Wasserstein Distance: Wasserstein GAN

1. **Wasserstein GAN (WGAN)**: WGANs model the class of 1-Lipschitz continuous functions using neural networks, denoted by f_ϕ , with the goal of

$$\min_{\theta} \max_{f_\phi} \left(\int f_\phi(\vec{x}) d\mathbb{P}_n(\vec{x}) - \int f_\phi(\vec{x}) d\mathbb{P}_\theta(\vec{x}) \right). \quad (35)$$

- When the set $\{f_\phi \mid \phi \in \Phi\}$ perfectly approximates the set $\{f \mid \|f\|_L \leq 1\}$, Equation (35) equals to $\min_{\theta} W_1(\mathbb{P}_n, \mathbb{P}_\theta; d)$.
- The 1-Lipschitz continuity condition is sometimes relaxed to C -Lipschitz continuity for an arbitrary constant C . To enforce this, WGANs clip weights and biases in neural network layers during training.

p -Wasserstein Distance: Wasserstein Autoencoder

2. **Wasserstein Autoencoder (WAE)**: Tolstikhin et al. (2018) derived an alternative representation of the p -Wasserstein distance:

$$W_p(\mathbb{P}_n, \mathbb{P}_\theta; d) = \left(\inf_{\mathbb{Q}(\vec{z}|\vec{x}): \int q(\vec{z}|\vec{x})d\mathbb{P}_n(\vec{x}) = p(\vec{z})} \int d^p(\vec{x}, G_\theta(\vec{z})) d\mathbb{Q}(\vec{z}|\vec{x}) d\mathbb{P}_n(\vec{x}) \right)^{1/p} \quad (36)$$

- Based on this relation, WAEs introduce encoders $q_\phi(\vec{z}|\vec{x})$ and target

$$\theta^* \in \arg \min_{\theta} \left(\inf_{\phi \in \Phi(\mathbb{P}_n)} \int d^p(\vec{x}, G_\theta(\vec{z})) d\mathbb{Q}_\phi(\vec{z}|\vec{x}) d\mathbb{P}_n(\vec{x}) \right)^{1/p} \quad (37)$$

where $\Phi(\mathbb{P}_n) := \{\phi \mid \int q_\phi(\vec{z}|\vec{x})d\mathbb{P}_n(\vec{x}) = p(\vec{z})\}$.

- On the RHS, $q_\phi(\vec{z}|\vec{x})$ can be viewed as an encoder. The constraint in the infimum ensures that the marginal distribution of the posterior distributions matches the prior distributions.

p -Wasserstein Distance: Wasserstein Autoencoder

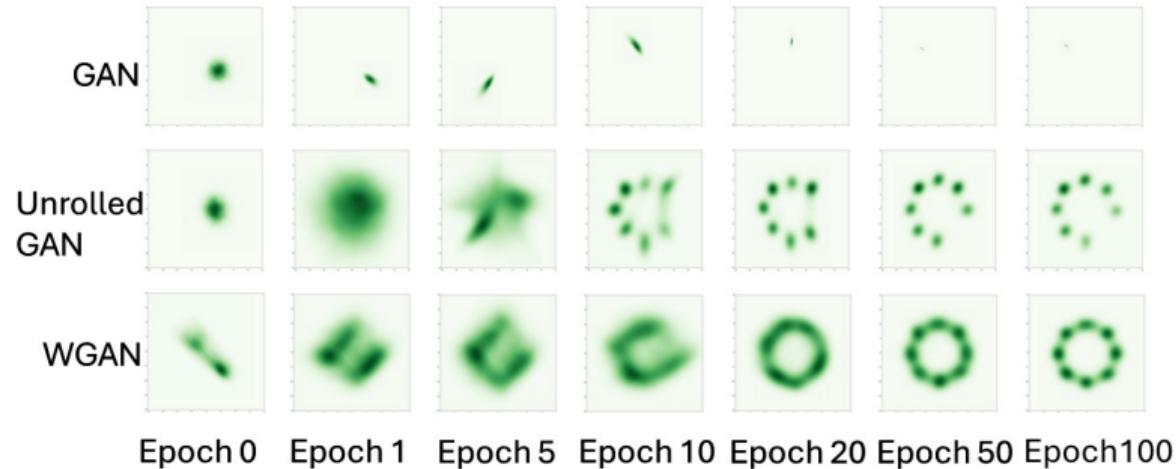
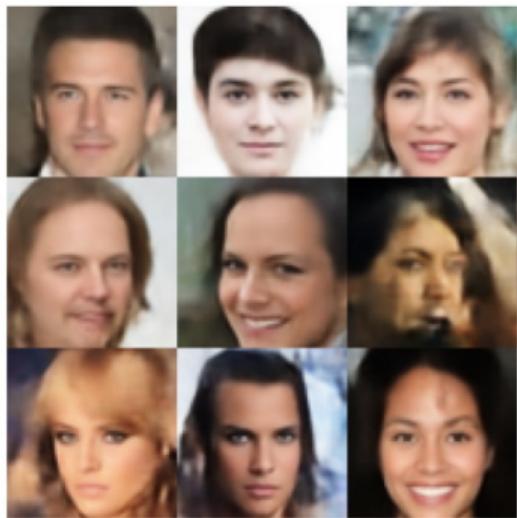
- In implementation, WAEs introduce a penalty term to enforce the constraint on ϕ . The loss can be expressed as:

$$\int d^P(\vec{x}, G_\theta(\vec{z}))d\mathbb{Q}_\phi(\vec{z}|\vec{x})d\mathbb{P}_n(\vec{x}) + \lambda \mathcal{D}_{\vec{Z}} \left(\int q_\phi(\vec{z}|\vec{x})d\mathbb{P}_n(\vec{x}), p(\vec{z}) \right) \quad (38)$$

where $\mathcal{D}_{\vec{Z}}$ indicates the statistical distance applied to the distributions of \vec{Z} . WAEs typically use JS divergence and MMD (Maximum Mean Discrepancy) as measures for $\mathcal{D}_{\vec{Z}}$.

- Compared with the loss of VAEs, the negative ELBO, the penalty term changes from matching $q_\phi(\vec{z}|\vec{x})$ directly with $p(\vec{z})$ to matching $\int q_\phi(\vec{z}|\vec{x})d\mathbb{P}_n(\vec{x})$ with $p(\vec{z})$.
- This difference in losses, motivated by theoretical results, may explain why WAEs often yield sharper and more plausible generative results compared to VAEs.

Generation Results: WAEs and WGANs



(a) Sharp generation results from WAEs

(b) Preventing mode collapse with WGANs

Images are edited from Arjovsky et al. (2017) and Tolstikhin et al. (2018).

Emergence of Fisher Divergence-based Methods



- IPM and Wasserstein distance-based methods have alleviated optimization issues; however, adversarial training is still practically difficult.
- Recent works have focused on score functions instead of densities, using estimated scores to generate data.

Recapping Fisher Divergence

- Fisher divergence (Johnson, 2004) is the expected difference between the (Stein) scores (Liu et al., 2016) of two distributions. It can be expressed as:

$$\text{FD}(p_n \parallel p_\theta) = \int \|\nabla_{\vec{x}} \log p_n(\vec{x}) - \nabla_{\vec{x}} \log p_\theta(\vec{x})\|^2 p_n(\vec{x}) d\vec{x}. \quad (39)$$

Fisher Divergence: Score Matching Estimation

1. **Score Matching Estimation:** Score matching estimation (Hyvärinen, 2005) was proposed targeting Fisher divergence in learning distributions.
- Let $S_\theta(\vec{x}) := \nabla_{\vec{x}} \log p_\theta(\vec{x})$. Then,

$$\text{FD}(p_n || p_\theta) = \int \left(\text{tr}(\nabla_{\vec{x}} S_\theta(\vec{x})) + \frac{1}{2} \|S_\theta(\vec{x})\|^2 \right) p_n(\vec{x}) d\vec{x} \quad (40)$$

up to a constant addition and sign-preserving multiplication. We assume that $S_\theta(\vec{x})p_n(\vec{x})$ vanishes at the boundary, e.g., $(x_1, \dots, x_{i-1}, \pm\infty, x_{i+1}, \dots, x_m)$.

Proof:

$$\begin{aligned} \text{FD}(p_n || p_\theta) &:= \int \|\nabla_{\vec{x}} \log p_n(\vec{x}) - S_\theta(\vec{x})\|^2 p_n(\vec{x}) d\vec{x} \\ &= C - 2 \int \left(S_\theta^T(\vec{x}) \nabla_{\vec{x}} \log p_n(\vec{x}) \right) p_n(\vec{x}) d\vec{x} + \int \|S_\theta(\vec{x})\|^2 p_n(\vec{x}) d\vec{x}. \end{aligned} \quad (41)$$

Here, $\int \left(S_\theta^T(\vec{x}) \nabla_{\vec{x}} \log p_n(\vec{x}) \right) p_n(\vec{x}) d\vec{x}$ equals $-\int \text{tr}(\nabla_{\vec{x}} S_\theta(\vec{x})) p_n(\vec{x}) d\vec{x}$.

Fisher Divergence: Score Matching Estimation

Proof (Cont.): Let $\vec{X}_{-i} := (X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_m)^T$. Then,

$$\begin{aligned} \int \left(S_\theta^T(\vec{x}) \nabla_{\vec{x}} \log p_n(\vec{x}) \right) p_n(\vec{x}) d\vec{x} &= \int S_\theta^T(\vec{x}) \nabla_{\vec{x}} p_n(\vec{x}) d\vec{x} \\ &= \sum_{i=1}^m \int \left(\int S_\theta(\vec{x})_i \frac{\partial}{\partial x_i} p_n(\vec{x}) dx_i \right) d\vec{x}_{-i}. \end{aligned} \tag{42}$$

Since $S_\theta(\vec{x})p_n(\vec{x})$ vanishes at the boundary, by partial integration, we have

$$\int S_\theta(\vec{x})_i \frac{\partial}{\partial x_i} p_n(\vec{x}) dx_i = - \int \left(\frac{\partial}{\partial x_i} S_\theta(\vec{x})_i \right) p_n(\vec{x}) dx_i. \tag{43}$$

Thus, $\text{FD}(p_n || p_\theta) = C + \int \left(2\text{tr}(\nabla_{\vec{x}} S_\theta(\vec{x})) + \|S_\theta(\vec{x})\|^2 \right) p_n(\vec{x}) d\vec{x}$, which concludes the proof.

Fisher Divergence: Sliced Score Matching

2. Sliced Score Matching:

- In the objective of score matching estimation, $\int \left(\text{tr}(\nabla_{\vec{x}} S_{\theta}(\vec{x})) + \frac{1}{2} \|S_{\theta}(\vec{x})\|^2 \right) p_n(\vec{x}) d\vec{x}$, the Hessian term poses another computational challenge.
- Sliced score matching (Song et al., 2020) targets sliced Fisher divergence (SFD),

$$\text{SFD}(p_n || p_{\theta}) := \int \left\| \vec{v}^T \nabla_{\vec{x}} \log p_n(\vec{x}) - \vec{v}^T \nabla_{\vec{x}} \log p_{\theta}(\vec{x}) \right\|^2 p_n(\vec{x}) p(\vec{v}) d\vec{x} d\vec{v}, \quad (44)$$

to overcome this limitation.

- The SFD is the average difference between randomly projected scores.

Fisher Divergence: Sliced Score Matching

- In a similar way used in score matching estimation,

$$\text{SFD}(p_n || p_\theta) = \int \left(\vec{v}^T \nabla_{\vec{x}} S_\theta(\vec{x}) \vec{v} + \frac{1}{2} (\vec{v}^T S_\theta(\vec{x}))^2 \right) p_n(\vec{x}) p(\vec{v}) d\vec{x} d\vec{v} \quad (45)$$

up to a constant addition and sign-preserving multiplication.

- By changing the target statistical distances from FD to SFD, the computational bottleneck shifts from computing $\text{tr}(\nabla_{\vec{x}} S_\theta(\vec{x}))$ to computing $\vec{v}^T \nabla_{\vec{x}} S_\theta(\vec{x}) = \nabla_{\vec{x}} (\vec{v}^T S_\theta(\vec{x}))$, which is numerically less demanding.
- When $p(\vec{v})$ is the multivariate standard Gaussian distribution, the equation $\int (\vec{v}^T S_\theta(\vec{x}))^2 d\vec{v} = \|S_\theta(\vec{x})\|^2$ holds, further reducing the computational cost.

Fisher Divergence: Noise Conditional Score Network

3. **Noise Conditional Score Network (NCSN)**: NCSNs (Song and Ermon, 2019) are score-based generative models that use estimated scores $S_\theta(\vec{x})$ to generate data.

- The key idea is to introduce Langevin dynamics in the sampling process. Langevin dynamics describes the stochastic movement of a fluid particle located at $\vec{X}(t)$:

$$m \frac{d^2 \vec{X}(t)}{dt^2} = -\nabla_{\vec{x}=\vec{X}(t)} U(\vec{x}) - \lambda \frac{d \vec{X}(t)}{dt} + \sqrt{2\lambda k_B T} \vec{B}(t), \quad (46)$$

where m is the mass, U is the potential functions, λ is the damping coefficient, k_B is the Boltzmann constant, T is the temperature, and $\vec{B}(t)$ represents the Brownian motion.

- In the overdamped case, where the inertial force is negligible, when $\lambda = 1$, we get

$$d\vec{X}(t) = -\nabla_{\vec{x}=\vec{X}(t)} U(\vec{x}) dt + \sqrt{2k_B T} d\vec{B}(t) \quad (47)$$

where $d\vec{B}(t) \sim N(0, dt I_m)$.¹⁰ Its stationary distribution is the Boltzmann distribution with energy $U/(k_B T)$, $p(\vec{x}(\infty)) \propto \exp(-U(\vec{x}(\infty))/(k_B T))$.

¹⁰This is a special case of the Itô drift-diffusion process.

Fisher Divergence: Noise Conditional Score Network

- By substituting $U(\vec{x}) = -\log p_n(\vec{x})$ and setting $T = 1/k_B$, we obtain:

$$d\vec{X}(t) = \nabla_{\vec{x}} \log p_n(\vec{x}) dt + \sqrt{2dt} \vec{\mathcal{E}}(t), \quad (48)$$

where $\vec{\mathcal{E}}(t) \sim N(0, I_m)$, and the corresponding stationary distribution is $p_n(\vec{x})$.

- The discrete approximation with $dt = \eta/2$ and $S_{\theta^*}(\vec{x})$ results in the following iterative sampling process:

$$\vec{X}(t) = \vec{X}(t-1) + (\eta/2) S_{\theta^*}(\vec{X}(t-1)) + \sqrt{\eta} \vec{\mathcal{E}}(t), \quad (49)$$

where $\vec{X}(T)$ approximately follows $p_n(\vec{x})$.

- Since the initial points are likely to lie in low-density regions, NCSNs employ the denoising score matching method (Vincent, 2011). They add noise to the data, $\vec{X} + \sigma \vec{\mathcal{E}}$, learn its score $S_\theta(\vec{x}; \sigma)$, and use $S_\theta(\vec{x}; \sigma)$ with a sufficiently small σ for effective sampling.

Summary



- We have reviewed recent developments in deep generative models, with a particular focus on targeted statistical distances.
- Advanced topics include:
 - ➊ Introducing new statistical distances,
 - ➋ Theoretical analysis of estimation and approximation errors,
 - ➌ Development of statistical models tailored to specific data structures, such as temporal or multi-modal data.



Thank You



MICHIGAN STATE
U N I V E R S I T Y

References I

- Arjovsky, M., Chintala, S., and Bottou, L. (2017). Wasserstein generative adversarial networks. In *International conference on machine learning*, pages 214–223. PMLR.
- Bengio, Y., Lamblin, P., Popovici, D., and Larochelle, H. (2006). Greedy layer-wise training of deep networks. *Advances in neural information processing systems*, 19.
- Bickel, P. J. and Freedman, D. A. (1981). Some asymptotic theory for the bootstrap. *The annals of statistics*, 9(6):1196–1217.
- Bishop, C. M. (2006). Pattern recognition and machine learning. *Springer google schola*, 2:1122–1128.
- Deng, L. (2012). The mnist database of handwritten digit images for machine learning research [best of the web]. *IEEE signal processing magazine*, 29(6):141–142.
- Freedman, D. A. (1981). Bootstrapping regression models. *The annals of statistics*, 9(6):1218–1228.

References II

- Georghiades, A. S., Belhumeur, P. N., and Kriegman, D. J. (2001). From few to many: Illumination cone models for face recognition under variable lighting and pose. *IEEE transactions on pattern analysis and machine intelligence*, 23(6):643–660.
- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative adversarial nets. *Advances in neural information processing systems*, 27.
- Gross, R., Matthews, I., Cohn, J., Kanade, T., and Baker, S. (2010). Multi-pie. *Image and vision computing*, 28(5):807–813.
- Gulrajani, I., Ahmed, F., Arjovsky, M., Dumoulin, V., and Courville, A. C. (2017). Improved training of wasserstein gans. *Advances in neural information processing systems*, 30.
- Hiriart-Urruty, J.-B. and Lemaréchal, C. (2004). *Fundamentals of convex analysis*. Springer Science & Business Media.

References III

- Ho, J., Jain, A., and Abbeel, P. (2020). Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851.
- Hyvärinen, A. (2005). Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 6(4).
- Johnson, O. (2004). *Information theory and the central limit theorem*. World Scientific.
- Kantorovich, L. V. (1960). Mathematical methods of organizing and planning production. *Management science*, 6(4):366–422.
- Kim, J.-H., Zhang, Y., Han, K., Wen, Z., Choi, M., and Liu, Z. (2021). Representation learning of resting state fmri with variational autoencoder. *NeuroImage*, 241:118423.
- Kingma, D. P. and Welling, M. (2014). Auto-encoding variational bayes. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

References IV

- Kramer, M. A. (1991). Nonlinear principal component analysis using autoassociative neural networks. *AIChE journal*, 37(2):233–243.
- Li, C.-L., Chang, W.-C., Cheng, Y., Yang, Y., and Póczos, B. (2017). Mmd gan: Towards deeper understanding of moment matching network. *Advances in neural information processing systems*, 30.
- Li, T., Tian, Y., Li, H., Deng, M., and He, K. (2024). Autoregressive image generation without vector quantization. *arXiv preprint arXiv:2406.11838*.
- Li, Y., Swersky, K., and Zemel, R. (2015). Generative moment matching networks. In *International conference on machine learning*, pages 1718–1727. PMLR.
- Liu, Q., Lee, J., and Jordan, M. (2016). A kernelized stein discrepancy for goodness-of-fit tests. In *International conference on machine learning*, pages 276–284. PMLR.
- Mallows, C. L. (1972). A note on asymptotic joint normality. *The Annals of Mathematical Statistics*, pages 508–515.

References V

- Metz, L., Poole, B., Pfau, D., and Sohl-Dickstein, J. (2017). Unrolled generative adversarial networks. In *International Conference on Learning Representations*.
- Monge, G. (1781). Mémoire sur la théorie des déblais et des remblais. *Mem. Math. Phys. Acad. Royale Sci.*, pages 666–704.
- Mroueh, Y., Li, C.-L., Sercu, T., Raj, A., and Cheng, Y. (2017). Sobolev gan. *arXiv preprint arXiv:1711.04894*.
- Müller, A. (1997). Integral probability metrics and their generating classes of functions. *Advances in applied probability*, 29(2):429–443.
- Nowozin, S., Cseke, B., and Tomioka, R. (2016). f-gan: Training generative neural samplers using variational divergence minimization. In *Advances in neural information processing systems*, pages 271–279.
- Plaut, E. (2018). From principal subspaces to principal components with linear autoencoders. *arXiv preprint arXiv:1804.10253*.

References VI

- Radford, A. (2015). Unsupervised representation learning with deep convolutional generative adversarial networks. *arXiv preprint arXiv:1511.06434*.
- Radford, A. (2018). Improving language understanding by generative pre-training.
- Ramesh, A., Dhariwal, P., Nichol, A., Chu, C., and Chen, M. (2022). Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*, 1(2):3.
- Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., Chen, M., and Sutskever, I. (2021). Zero-shot text-to-image generation. In *International conference on machine learning*, pages 8821–8831. Pmlr.
- Rényi, A. (1961). On measures of entropy and information. In *Proceedings of the fourth Berkeley symposium on mathematical statistics and probability, volume 1: contributions to the theory of statistics*, volume 4, pages 547–562. University of California Press.
- Santambrogio, F. (2015). Optimal transport for applied mathematicians. *Birkhäuser, NY*, 55(58-63):94.

References VII

- Song, Y. and Ermon, S. (2019). Generative modeling by estimating gradients of the data distribution. *Advances in neural information processing systems*, 32.
- Song, Y., Garg, S., Shi, J., and Ermon, S. (2020). Sliced score matching: A scalable approach to density and score estimation. In *Uncertainty in Artificial Intelligence*, pages 574–584. PMLR.
- Teh, Y. W., Welling, M., Osindero, S., and Hinton, G. E. (2003). Energy-based models for sparse overcomplete representations. *Journal of Machine Learning Research*, 4(Dec):1235–1260.
- Tolstikhin, I., Bousquet, O., Gelly, S., and Schölkopf, B. (2018). Wasserstein auto-encoders. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Uppal, A., Singh, S., and Póczos, B. (2019). Nonparametric density estimation & convergence rates for gans under besov ipm losses. *Advances in neural information processing systems*, 32.

References VIII

- Villani, C. (2021). *Topics in optimal transportation*, volume 58. American Mathematical Soc.
- Villani, C. et al. (2009). *Optimal transport: old and new*, volume 338. Springer.
- Vincent, P. (2011). A connection between score matching and denoising autoencoders. *Neural computation*, 23(7):1661–1674.
- Wolterink, J. M., Dinkla, A. M., Savenije, M. H., Seevinck, P. R., van den Berg, C. A., and Išgum, I. (2017). Deep mr to ct synthesis using unpaired data. In *International workshop on simulation and synthesis in medical imaging*, pages 14–23. Springer.